



PERCONA

Databases run better with Percona

October 2024



PERCONA

Building Out Active-Active Replication Clusters,

Architectural Considerations



Robert Bernier

Robert.bernier@percona.com

Senior PostgreSQL Consultant

Percona

About This Talk

- About This Talk
- Active-Active Replication
 - Understanding Active-Active
 - Active-Active: 2 Node example
- Architectural Types
 - Daisy Chain
 - Star
 - Mesh
 - Xmas Tree
- Administering The Logical Replication Cluster
 - About
 - Best Practices
 - Caveat
 - Conflict Resolution
 - Monitoring

Active-Active



Understanding Active-Active

● Key Concerns

- DML conflicts when the subscribed data conflicts with data already present in the relation.
- The "echo effect" where data, initially propagated from the published table, returns to its original source host server as a DML operation.

● Mitigation Techniques

- Maintaining consistent table definitions i.e. table columns and their constraints between PUBLISHED and SUBSCRIBED tables.
- Filtering undesired rows when creating the PUBLICATION using the WHERE clause
- Setting the ORIGIN parameter to "none" at SUBSCRIPTION creation.

Active-Active: 2 Node example

Step 1: Create Environment

```
-- common to both primary nodes  
create role active_active with login replication password 'active';  
create database db01;
```

Active-Active: 2 Node example

Step 1 cont'd

```
-- pg1
\c db01
drop table if exists t1;
create table t1(
    id uuid default gen_random_uuid()
    ,comments text default 'pg1'
    ,t_stamp timestamptz default clock_timestamp()
    ,PRIMARY KEY (t_stamp,id)
);

grant all on all tables in schema public to active_active;
```

```
-- pg2
\c db01
drop table if exists t1;
create table t1(
    id uuid default gen_random_uuid()
    ,comments text default 'pg2'
    ,t_stamp timestamptz default clock_timestamp()
    ,PRIMARY KEY (t_stamp,id)
);

grant all on all tables in schema public to active_active;
```


Active-Active: 2 Node example

Step 2: Start Replication

```
-- PUBLICATION
```

```
-- pg1  
create publication pg1 for table t1;
```

```
-- pg2  
create publication pg2 for table t1;
```

```
-- SUBSCRIPTION
```

```
-- pg1  
create subscription pg1  
  connection 'host=pg2 port=5432 dbname=db01  
  user=active_active password=active'  
  publication pg2  
  with (origin=none, copy_data=false);
```

```
-- pg2  
create subscription pg2  
  connection 'host=pg1 port=5432 dbname=db01  
  user=active_active password=active'  
  publication pg1  
  with (origin=none, copy_data=true);
```

Active-Active: 2 Node example

Step 3: Validation

```
-- pg1
insert into t1 values
  (default,default,default),
  (default,default,default),
  (default,default,default);
```

```
-- pg2
insert into t1 values
  (default,default,default),
  (default,default,default),
  (default,default,default);
```

Active-Active: Caveat

- In order that logical replication can function all tuples (records) must be unique in a table i.e. possess a PRIMARY, or UNIQUE, key..
- Where primary or unique keys do not exist in a table, the SQL command `ALTER TABLE ... REPLICA IDENTITY ...` can add the necessary information to the WAL making logical replication possible.

Architectural Types

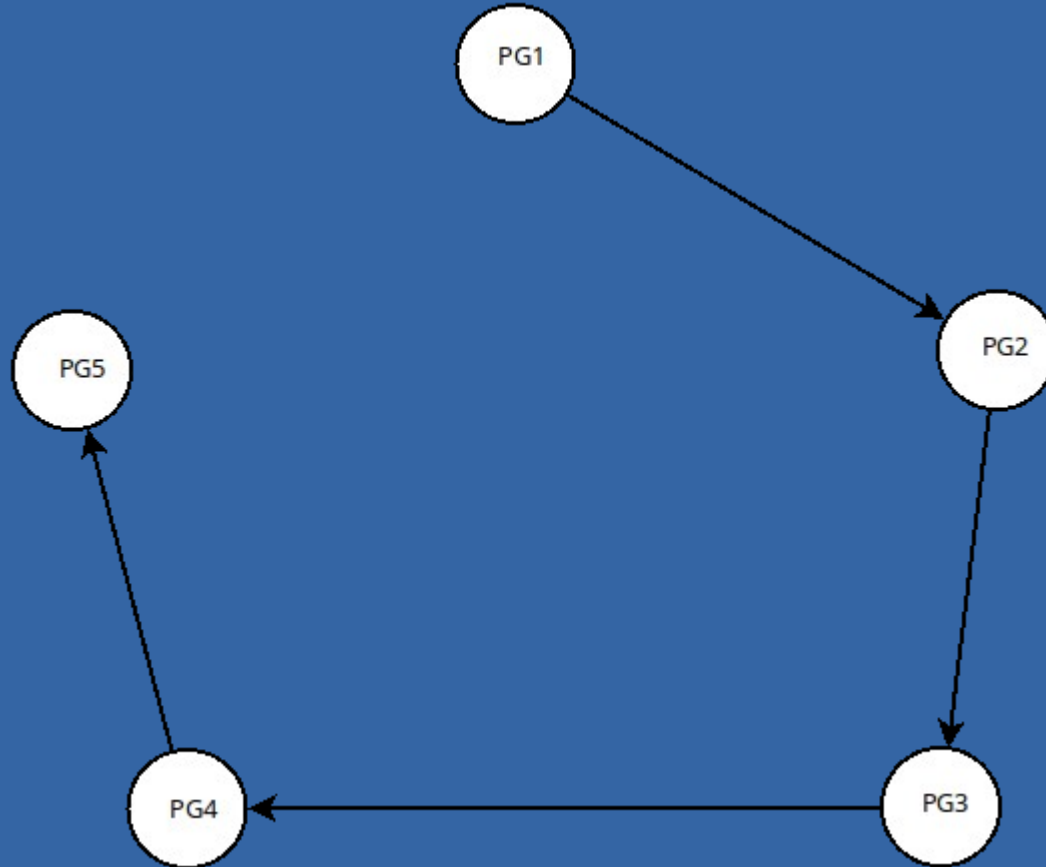


Architectural Types

- Daisy Chain
- Star
- Mesh

Daisy Chain

(1 way-replication)



Daisy Chain

Step 1: Create Table and Publication

```
for u in 1 2 3 4 5
do
    echo "===== pg$u ====="
    psql -q "host=pg$u dbname=db01 user=postgres password=postgres" <<_eof_
        drop publication if exists pg$u;
        drop table if exists t1;
        create table t1(
            id uuid default gen_random_uuid()
            ,comments text default 'pg$u-openai'
            ,t_stamp timestamptz default clock_timestamp()
            ,PRIMARY KEY (t_stamp,id)
        );
        grant all on all tables in schema public to active_active;
        create publication pg$u for table t1;
_eof_
done
```

Daisy Chain

Step 2: Create the Subscription

```
for u in 2 3 4 5
do
    let x=$u-1
    echo "===== pg$x -> pg$u ====="
    psql -q "host=pg$u dbname=db01 user=postgres password=postgres" <<_eof_
        create subscription pg$x
            connection 'host=pg$x port=5432 dbname=db01 user=active_active
password=active'
            publication pg$x
            with (origin=any, copy_data=false, slot_name=pg$u);
_eof_
done
```


Daisy Chain

Step 3: Insert Records

```
for u in 1 2 3 4 5
do
    echo "===== pg${u} ====="
    psql -q "host=pg${u} dbname=db01 user=postgres password=postgres" <<_eof_
        insert into t1 values
            (default,default,default),
            (default,default,default);
    _eof_
    sleep 1s
done
```

Daisy Chain

Step 4: Query Records

```
for u in 1 2 3 4 5
do
    echo "===== pg${u} ====="
    psql -q "host=pg${u} dbname=db01 user=postgres password=postgres" <<_eof_
    select * from t1 order by 3;
_eof_
done
```

Daisy Chain

Step 4 cont'd:

- Advantages ?
- Disadvantages?

```
===== pg1 =====
      id | comments | t_stamp
-----+-----+-----
8d32985a-b78e-421f-a679-1090275a0c38 | pg1-openai | 2024-06-25 21:17:11.958845+00
868732b3-5e1a-46f8-8606-8c418f79240b | pg1-openai | 2024-06-25 21:17:11.958971+00
(2 rows)

===== pg2 =====
      id | comments | t_stamp
-----+-----+-----
8d32985a-b78e-421f-a679-1090275a0c38 | pg1-openai | 2024-06-25 21:17:11.958845+00
868732b3-5e1a-46f8-8606-8c418f79240b | pg1-openai | 2024-06-25 21:17:11.958971+00
266ce44c-a896-4f97-b2ce-cee4ef4fb15d | pg2-openai | 2024-06-25 21:17:13.000763+00
8075f71c-5af5-45db-9db6-5e0e679e4ada | pg2-openai | 2024-06-25 21:17:13.000832+00
(4 rows)

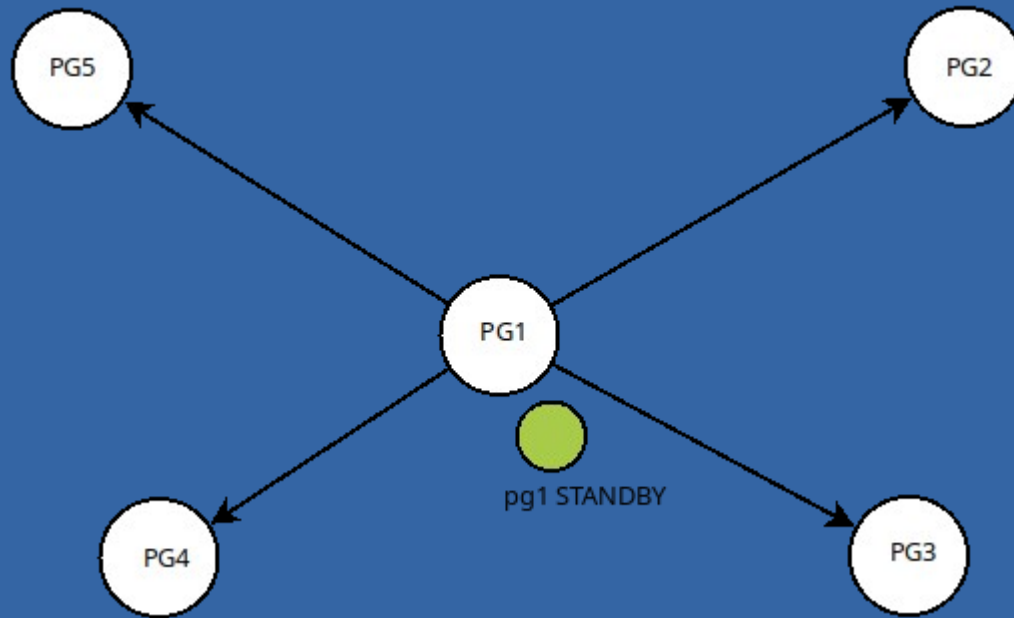
===== pg3 =====
      id | comments | t_stamp
-----+-----+-----
8d32985a-b78e-421f-a679-1090275a0c38 | pg1-openai | 2024-06-25 21:17:11.958845+00
868732b3-5e1a-46f8-8606-8c418f79240b | pg1-openai | 2024-06-25 21:17:11.958971+00
266ce44c-a896-4f97-b2ce-cee4ef4fb15d | pg2-openai | 2024-06-25 21:17:13.000763+00
8075f71c-5af5-45db-9db6-5e0e679e4ada | pg2-openai | 2024-06-25 21:17:13.000832+00
e060482e-1e4d-48d9-88ec-733d738b741c | pg3-openai | 2024-06-25 21:17:14.041163+00
d3414363-bb00-489f-94b8-e37908a729f4 | pg3-openai | 2024-06-25 21:17:14.041231+00
(6 rows)

===== pg4 =====
      id | comments | t_stamp
-----+-----+-----
8d32985a-b78e-421f-a679-1090275a0c38 | pg1-openai | 2024-06-25 21:17:11.958845+00
868732b3-5e1a-46f8-8606-8c418f79240b | pg1-openai | 2024-06-25 21:17:11.958971+00
266ce44c-a896-4f97-b2ce-cee4ef4fb15d | pg2-openai | 2024-06-25 21:17:13.000763+00
8075f71c-5af5-45db-9db6-5e0e679e4ada | pg2-openai | 2024-06-25 21:17:13.000832+00
e060482e-1e4d-48d9-88ec-733d738b741c | pg3-openai | 2024-06-25 21:17:14.041163+00
d3414363-bb00-489f-94b8-e37908a729f4 | pg3-openai | 2024-06-25 21:17:14.041231+00
6225fc7f-038d-4001-a8ad-e2f9c544bee2 | pg4-openai | 2024-06-25 21:17:15.081426+00
3bbc856c-65e0-4789-a106-b3a2f29097dc | pg4-openai | 2024-06-25 21:17:15.081499+00
(8 rows)

===== pg5 =====
      id | comments | t_stamp
-----+-----+-----
8d32985a-b78e-421f-a679-1090275a0c38 | pg1-openai | 2024-06-25 21:17:11.958845+00
868732b3-5e1a-46f8-8606-8c418f79240b | pg1-openai | 2024-06-25 21:17:11.958971+00
266ce44c-a896-4f97-b2ce-cee4ef4fb15d | pg2-openai | 2024-06-25 21:17:13.000763+00
8075f71c-5af5-45db-9db6-5e0e679e4ada | pg2-openai | 2024-06-25 21:17:13.000832+00
e060482e-1e4d-48d9-88ec-733d738b741c | pg3-openai | 2024-06-25 21:17:14.041163+00
d3414363-bb00-489f-94b8-e37908a729f4 | pg3-openai | 2024-06-25 21:17:14.041231+00
6225fc7f-038d-4001-a8ad-e2f9c544bee2 | pg4-openai | 2024-06-25 21:17:15.081426+00
3bbc856c-65e0-4789-a106-b3a2f29097dc | pg4-openai | 2024-06-25 21:17:15.081499+00
a9690074-2eff-423d-bec6-8a81c5043f27 | pg5-openai | 2024-06-25 21:17:16.121258+00
169d3989-c951-480d-8d21-6fc24dbde30d | pg5-openai | 2024-06-25 21:17:16.121351+00
(10 rows)
```

Star Topology

(1 way-replication)



Star Topology

Step 1: Create Table and Publication on Primary

```
psql "host=pg1 dbname=db01 user=postgres password=postgres" <<_eof_
drop publication if exists pg1;
drop table if exists t1;
create table t1(
    id uuid default gen_random_uuid()
    ,comments text default 'pg1-openai'
    ,t_stamp timestamptz default clock_timestamp()
    ,PRIMARY KEY (t_stamp,id)
);
create publication pg1 for table t1;
grant all on all tables in schema public to active_active;
_eof_
```

Star Topology

Step 1 cont'd: Create Table on REPLICA(s)

```
psql "host=pg1 dbname=db01 user=postgres password=postgres" <<_eof_  
drop publication if exists pg1;  
drop table if exists t1;  
create table t1(  
    id uuid default gen_random_uuid()  
    ,comments text default 'pg1-openai'  
    ,t_stamp timestamptz default clock_timestamp()  
    ,PRIMARY KEY (t_stamp,id)  
);  
_eof_
```

Star Topology

Step 2: Create Subscription

```
for u in 2 3 4 5
do
    echo "===== pg1 -> pg$u ====="
    psql -q "host=pg$u dbname=db01 user=postgres password=postgres" <<_eof_
    create subscription pg1
    connection 'host=pg1 port=5432 dbname=db01 user=active_active password=active'
    publication pg1
    with (origin=none, copy_data=false, slot_name=pg$u);
_eof_
echo "===== pg$u ====="
psql -q "host=pg$u dbname=db01 user=postgres password=postgres" <<_eof_
select subname, subenabled, subconninfo, subslotname from
    pg_subscription;
_eof_
done
```

Star Topology

step 2 cont'd:

```
===== pg1 -> pg2 =====
WARNING: publication "pg2" does not exist on the publisher
NOTICE: created replication slot "pg2" on publisher
===== pg2 =====
  subname | subenabled |                               subconninfo                               | subslotname
-----+-----+-----+-----+-----+-----+-----+-----+-----+
  pg1     | t          | host=pg1 port=5432 dbname=db01 user=active_active password=active | pg2
(1 row)

===== pg1 -> pg3 =====
WARNING: publication "pg3" does not exist on the publisher
NOTICE: created replication slot "pg3" on publisher
===== pg3 =====
  subname | subenabled |                               subconninfo                               | subslotname
-----+-----+-----+-----+-----+-----+-----+-----+
  pg1     | t          | host=pg1 port=5432 dbname=db01 user=active_active password=active | pg3
(1 row)

===== pg1 -> pg4 =====
WARNING: publication "pg4" does not exist on the publisher
NOTICE: created replication slot "pg4" on publisher
===== pg4 =====
  subname | subenabled |                               subconninfo                               | subslotname
-----+-----+-----+-----+-----+-----+-----+-----+
  pg1     | t          | host=pg1 port=5432 dbname=db01 user=active_active password=active | pg4
(1 row)

===== pg1 -> pg5 =====
WARNING: publication "pg5" does not exist on the publisher
NOTICE: created replication slot "pg5" on publisher
===== pg5 =====
  subname | subenabled |                               subconninfo                               | subslotname
-----+-----+-----+-----+-----+-----+-----+-----+
  pg1     | t          | host=pg1 port=5432 dbname=db01 user=active_active password=active | pg5
(1 row)
```


Star Topology

Step 3: Insert Records

```
for u in 1 2 3 4 5
do
    echo "===== pg${u} ====="
    psql -q "host=pg${u} dbname=db01 user=postgres password=postgres" <<_eof_
        insert into t1 values
            (default,default,default),
            (default,default,default);
    _eof_
done
```

Star Topology

Step 4: Query Records

```
for u in 1 2 3 4 5
do
    echo "===== pg${u} ====="
    psql -q "host=pg${u} dbname=db01 user=postgres password=postgres" <<_eof_
    select * from t1 order by 3;
_eof_
done
```

Star Topology

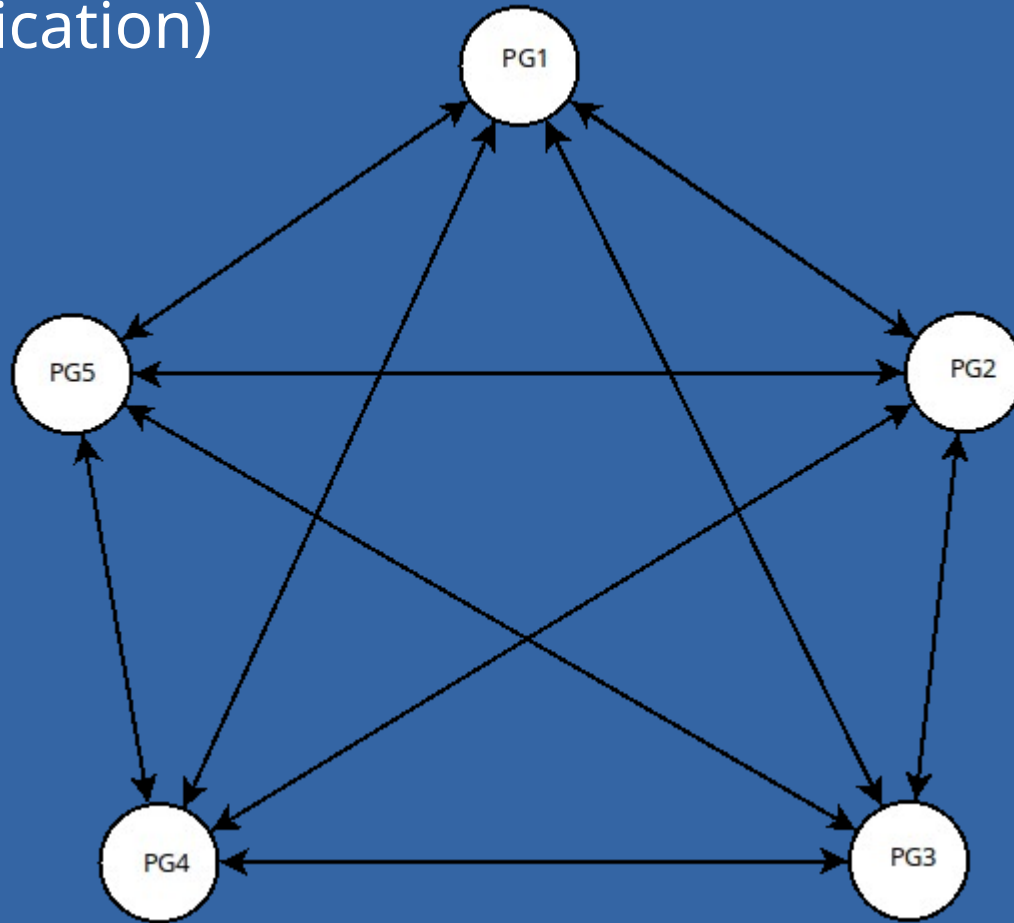
step 4 cont'd:

- Advantages ?
- Disadvantages?

===== pg1 =====			
id	comments	t_stamp	
8ede70f3-8409-4783-b03a-f05add1899e2	pg1-openai	2024-06-26	19:10:26.834392+00
9ef5d0db-4a93-4886-b489-4a22f0950068	pg1-openai	2024-06-26	19:10:26.834529+00
(2 rows)			
===== pg2 =====			
id	comments	t_stamp	
8ede70f3-8409-4783-b03a-f05add1899e2	pg1-openai	2024-06-26	19:10:26.834392+00
9ef5d0db-4a93-4886-b489-4a22f0950068	pg1-openai	2024-06-26	19:10:26.834529+00
03dd34bc-2212-4837-a6bd-ebd9d10e0a51	pg2-openai	2024-06-26	19:10:26.874537+00
69d982e2-445d-4add-86e6-5d4ee6ccf753	pg2-openai	2024-06-26	19:10:26.874612+00
(4 rows)			
===== pg3 =====			
id	comments	t_stamp	
8ede70f3-8409-4783-b03a-f05add1899e2	pg1-openai	2024-06-26	19:10:26.834392+00
9ef5d0db-4a93-4886-b489-4a22f0950068	pg1-openai	2024-06-26	19:10:26.834529+00
c4b7f0f1-eba2-4c16-ab69-930dd7ed701a	pg3-openai	2024-06-26	19:10:26.910106+00
ece3e682-9105-4aa7-ba43-7b75867285f0	pg3-openai	2024-06-26	19:10:26.910186+00
(4 rows)			
===== pg4 =====			
id	comments	t_stamp	
8ede70f3-8409-4783-b03a-f05add1899e2	pg1-openai	2024-06-26	19:10:26.834392+00
9ef5d0db-4a93-4886-b489-4a22f0950068	pg1-openai	2024-06-26	19:10:26.834529+00
6ed4c0e6-bc00-48b6-9817-3e8d0cee3e13	pg4-openai	2024-06-26	19:10:26.943257+00
4894ce41-2930-4663-b1c2-76a530f485e0	pg4-openai	2024-06-26	19:10:26.943337+00
(4 rows)			
===== pg5 =====			
id	comments	t_stamp	
8ede70f3-8409-4783-b03a-f05add1899e2	pg1-openai	2024-06-26	19:10:26.834392+00
9ef5d0db-4a93-4886-b489-4a22f0950068	pg1-openai	2024-06-26	19:10:26.834529+00
5f68691c-1a3e-4d35-9a0b-a53ee8a2eae1	pg5-openai	2024-06-26	19:10:26.978153+00
cbba47f3-a0f7-4735-86b2-96f01cbe7307	pg5-openai	2024-06-26	19:10:26.978228+00
(4 rows)			

Mesh

(2 way-replication)



Mesh

Step 1: Create Table and Publication on nodes

```
for u in 1 2 3 4 5
do
echo "===== pg${u} ====="
psql "host=pg${u} dbname=db01 user=postgres password=postgres" <<_eof_
    drop publication if exists pg${u};
    drop table if exists t1;
    create table t1(
        id uuid default gen_random_uuid()
        ,comments text default 'pg${u}-openai'
        ,t_stamp timestamptz default clock_timestamp()
        ,PRIMARY KEY (t_stamp,id)
    );
    grant all on all tables in schema public to active_active;
    create publication pg${u} for table t1;
_eof_
done
```

Mesh

Step 2: Create Subscription

```
for u in 1 2 3 4 5
do
    for v in 1 2 3 4 5
    do
        if [ $u -ne $v ]
        then
            echo "===== pg$u -> pg$v ====="
            psql -q "host=pg$u dbname=db01 user=postgres password=postgres" <<_eof_
                create subscription pg$v
                    connection 'host=pg$v port=5432 dbname=db01 user=active_active
password=active'
                    publication pg$v
                    with (origin=none, copy_data=false, slot_name=pg$u);
            _eof_
        fi
    done
done
```

Mesh

Step 3: Insert Records

```
for u in 1 2 3 4 5
do
    echo "===== pg${u} ====="
    psql -q "host=pg${u} dbname=db01 user=postgres password=postgres" <<_eof_
        insert into t1 values
            (default,default,default),
            (default,default,default);
    _eof_
    sleep 1s
done
```

Mesh

Step 4: Query Records

```
for u in 1 2 3 4 5
do
    echo "===== pg${u} ====="
    psql -q "host=pg${u} dbname=db01 user=postgres password=postgres" <<_eof_
        select * from t1 order by 3;
_eof_
done
```


Mesh

step 4 cont'd:

```
===== pg1 =====
-----
id | comments | t_stamp
-----
5fab5f67-6cea-465c-9eb7-e46787ab84b0 | pg1-openai | 2024-06-26 17:12:01.9294+00
08537516-bd1a-4849-b2da-8eec727529de | pg1-openai | 2024-06-26 17:12:01.929519+00
72e9f04e-74fb-4521-9824-971bf5093340 | pg2-openai | 2024-06-26 17:12:02.970595+00
3f855892-899d-4da4-8ae1-1e4b14025c9d | pg2-openai | 2024-06-26 17:12:02.97066+00
12f8318d-4116-49ba-9643-217429637c30 | pg3-openai | 2024-06-26 17:12:04.013467+00
70655e17-0d0c-4c16-a4ba-15c259dbbe39 | pg3-openai | 2024-06-26 17:12:04.013531+00
61d5fa4b-2efd-44d9-8c9b-fb0d426c0b62 | pg4-openai | 2024-06-26 17:12:05.054036+00
0c70375a-606a-4d85-908d-b4e5dd313293 | pg4-openai | 2024-06-26 17:12:05.054114+00
3e020797-7a64-42b9-91df-4ea3f7011c7c | pg5-openai | 2024-06-26 17:12:06.094513+00
c856550c-975b-43a1-84e5-61194a76ffdd | pg5-openai | 2024-06-26 17:12:06.094584+00
(10 rows)
```

```
===== pg2 =====
-----
id | comments | t_stamp
-----
5fab5f67-6cea-465c-9eb7-e46787ab84b0 | pg1-openai | 2024-06-26 17:12:01.9294+00
08537516-bd1a-4849-b2da-8eec727529de | pg1-openai | 2024-06-26 17:12:01.929519+00
72e9f04e-74fb-4521-9824-971bf5093340 | pg2-openai | 2024-06-26 17:12:02.970595+00
3f855892-899d-4da4-8ae1-1e4b14025c9d | pg2-openai | 2024-06-26 17:12:02.97066+00
12f8318d-4116-49ba-9643-217429637c30 | pg3-openai | 2024-06-26 17:12:04.013467+00
70655e17-0d0c-4c16-a4ba-15c259dbbe39 | pg3-openai | 2024-06-26 17:12:04.013531+00
61d5fa4b-2efd-44d9-8c9b-fb0d426c0b62 | pg4-openai | 2024-06-26 17:12:05.054036+00
0c70375a-606a-4d85-908d-b4e5dd313293 | pg4-openai | 2024-06-26 17:12:05.054114+00
3e020797-7a64-42b9-91df-4ea3f7011c7c | pg5-openai | 2024-06-26 17:12:06.094513+00
c856550c-975b-43a1-84e5-61194a76ffdd | pg5-openai | 2024-06-26 17:12:06.094584+00
(10 rows)
```

```
===== pg3 =====
-----
id | comments | t_stamp
-----
5fab5f67-6cea-465c-9eb7-e46787ab84b0 | pg1-openai | 2024-06-26 17:12:01.9294+00
08537516-bd1a-4849-b2da-8eec727529de | pg1-openai | 2024-06-26 17:12:01.929519+00
72e9f04e-74fb-4521-9824-971bf5093340 | pg2-openai | 2024-06-26 17:12:02.970595+00
3f855892-899d-4da4-8ae1-1e4b14025c9d | pg2-openai | 2024-06-26 17:12:02.97066+00
12f8318d-4116-49ba-9643-217429637c30 | pg3-openai | 2024-06-26 17:12:04.013467+00
70655e17-0d0c-4c16-a4ba-15c259dbbe39 | pg3-openai | 2024-06-26 17:12:04.013531+00
61d5fa4b-2efd-44d9-8c9b-fb0d426c0b62 | pg4-openai | 2024-06-26 17:12:05.054036+00
0c70375a-606a-4d85-908d-b4e5dd313293 | pg4-openai | 2024-06-26 17:12:05.054114+00
3e020797-7a64-42b9-91df-4ea3f7011c7c | pg5-openai | 2024-06-26 17:12:06.094513+00
c856550c-975b-43a1-84e5-61194a76ffdd | pg5-openai | 2024-06-26 17:12:06.094584+00
(10 rows)
```

```
===== pg4 =====
-----
id | comments | t_stamp
-----
5fab5f67-6cea-465c-9eb7-e46787ab84b0 | pg1-openai | 2024-06-26 17:12:01.9294+00
08537516-bd1a-4849-b2da-8eec727529de | pg1-openai | 2024-06-26 17:12:01.929519+00
```

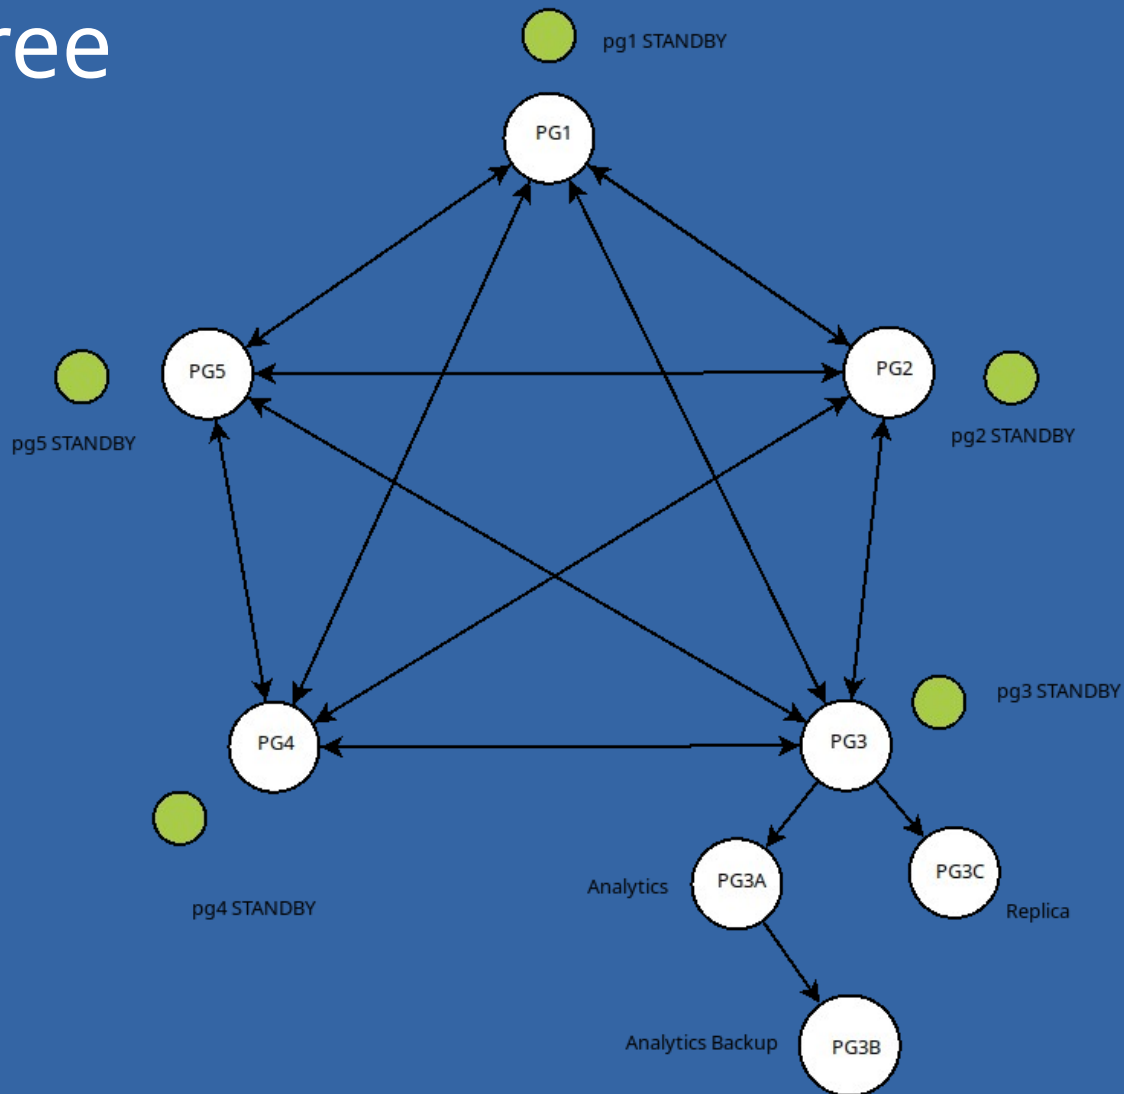
Mesh

- **Advantages:** Not only can multiple DML operations be executed on each host but one can failover to any one of the PRIMARY read-write nodes without downtime or SUBSCRIPTION configuration updates.
- **Disadvantages:** Topology complexity and network overhead increases as the system scales.

Mesh

Number Of Nodes	Total Number Slots req'd
2	2
3	6
4	12
5	20
6	30

XMAS Tree



Questions?

https://github.com/rbernierZulu/pg_conf_Seattle-2024

Thank You!