

Working With PostgreSQL as a JSON Document Store

Robert Bernier

Robert.bernier@percona.com

Senior PostgreSQL Consultant

Percona



Howdy!

Who am I?

Howdy!

Who RU?

Make Sense of Your Database

About This Talk

- **Part I: An Overview of PostgreSQL JSON**
- **Part II: Scenarios**

Make Sense of Your Database

PART I: An Overview of PostgreSQL JSON

Some Definitions

- 1) DBMS
- 2) RDBMS
- 3) GPDBMS

PostgreSQL: General Purpose Database Management System

- Relational (structured)
 - row wise: OLTP
 - column wise: OLAP
 - data warehouse
 - analytics
- Unstructured
 - hash
 - hierarchical
 - document storage
- High Availability (HA)
 - logical backups
 - clustered, across nodes ex: replication
- High Performance (HP)
 - 1) clustered ex: horizontal scaling
 - sharding with data redundancies
- Embedded (application centric)
- Big Data (inter/intra data centres)
- Encrypted
 - session
 - data
 - at rest (encrypted discs)

About JSON

Unstructured data files often include text and multimedia content. Examples include e-mail messages, word processing documents, videos, photos, audio files, presentations, webpages, and many other kinds of business documents.

JSON (JavaScript Object Notation) is a lightweight data-interchange format. It is easy for humans to read and write. It is easy for machines to parse and generate. It is based on a subset of the JavaScript Programming Language Standard ECMA-262 3rd Edition - December 1999.

JSON is a text format that is completely language independent but uses conventions that are familiar to programmers of the C-family of languages, including C, C++, C#, Java, JavaScript, Perl, Python, and many others. These properties make JSON an ideal data-interchange language.

<http://www.ecma-international.org/publications/files/ECMA-ST/ECMA-404.pdf>

About the PostgreSQL Implementation of JSON

- 1) **Data Types**
- 2) **Indexes**
- 3) **JSON Path Type**
- 4) **Operators**
- 5) **Functions**

Data Types

```
create table mytable (  
  id integer primary key generated by default as identity,  
  mydoc json,  
  mydocb jsonb  
);
```

```
Table "public.mytable"  
Column | Type      | Collation | Nullable | Default  
-----+-----+-----+-----+-----  
id      | integer   |           | not null | generated by default as identity  
mydoc   | json      |           |          |  
mydocb  | jsonb     |           |          |  
Indexes:  
    "mytable_pkey" PRIMARY KEY, btree (id)
```

<https://www.postgresql.org/docs/current/datatype-json.html>

Example: Simple Scalar/Primitive Values

-- Primitive values can be numbers, quoted strings, true, false, or null

```
SELECT '5'::json;
```

-- Array of zero or more elements (elements need not be of same type)

```
SELECT '[1, 2, "foo", null]'::json;
```

-- Object containing pairs of keys and values

-- Note that object keys must always be quoted strings

```
SELECT '{"bar": "baz", "balance": 7.77, "active": false}'::json;
```

-- Arrays and objects can be nested arbitrarily

```
SELECT '{"foo": [true, "bar"], "tags": {"a": 1, "b": null}}'::json;
```

Example: JSONB Containment and Existence

-- Simple scalar/primitive values contain only the identical value:

```
SELECT '"foo"'::jsonb @> '"foo"'::jsonb;
```

-- The array on the right side is contained within the one on the left:

```
SELECT '[1, 2, 3]'::jsonb @> '[1, 3]'::jsonb;
```

-- Order of array elements is not significant, so this is also true:

```
SELECT '[1, 2, 3]'::jsonb @> '[3, 1]'::jsonb;
```

-- Duplicate array elements don't matter either:

```
SELECT '[1, 2, 3]'::jsonb @> '[1, 2, 2]'::jsonb;
```

-- The object with a single pair on the right side is contained

-- within the object on the left side:

```
SELECT '{"product": "PostgreSQL", "version": 9.4, "jsonb": true}'::jsonb @> '{"version": 9.4}'::jsonb;
```

-- The array on the right side is not considered contained within the

-- array on the left, even though a similar array is nested within it:

```
SELECT '[1, 2, [1, 3]]'::jsonb @> '[1, 3]'::jsonb; -- yields false
```

-- But with a layer of nesting, it is contained:

```
SELECT '[1, 2, [1, 3]]'::jsonb @> '[[1, 3]]'::jsonb;
```

-- Similarly, containment is not reported here:

```
SELECT '{"foo": {"bar": "baz"}}'::jsonb @> '{"bar": "baz"}'::jsonb; -- yields false
```

-- A top-level key and an empty object is contained:

```
SELECT '{"foo": {"bar": "baz"}}'::jsonb @> '{"foo": {}}'::jsonb;
```

Indexes

```
create table mytable (  
  id integer primary key generated by default as identity,  
  mydoc json,  
  mydocb jsonb  
);
```

-- only JSONB

```
create index on mytable using gin (mydocb);
```

Table "public.mytable"				
Column	Type	Collation	Nullable	Default
id	integer		not null	generated by default as identity
mydoc	json			
mydocb	jsonb			

Indexes:

```
"mytable_pkey" PRIMARY KEY, btree (id)  
"mytable_mydocb_idx" gin (mydocb)
```

Example

```
{  
  "guid": "9c36adc1-7fb5-4d5b-83b4-90356a46061a",  
  "name": "Angela Barton",  
  "is_active": true,  
  "company": "MagnaFone",  
  "address": "178 Howard Place, Gulf, Washington, 702",  
  "registered": "2009-11-07T08:53:22 +08:00",  
  "latitude": 19.793713,  
  "longitude": 86.513373,  
  "tags": [  
    "enim",  
    "aliquip",  
    "qui"  
  ]  
}
```

Example, cont'd

-
- Find documents in which the key "company" has value "Magnafone"
- Does the left JSON value contain the right
- JSON path/value entries at the top level?
-

```
tmp=# EXPLAIN
tmp=# SELECT mydocb->'guid', mydocb->'name'
tmp=# FROM mytable
tmp=# WHERE mydocb @> '{"company": "Magnafone"}';
QUERY PLAN
```

```
-----
Seq Scan on mytable (cost=0.00..20.63 rows=1 width=64)
  Filter: (mydocb @> '{"company": "Magnafone"}'::jsonb)
```

Example, cont'd

```
tmp=# create index on mytable using gin ((mydocb -> 'tags'));  
--
```

```
tmp=# EXPLAIN
```

```
tmp=# SELECT mydocb->'guid', mydocb->'name'
```

```
tmp=# FROM mytable
```

```
tmp=# WHERE mydocb @> '{"company": "Magnafone"}';
```

```
QUERY PLAN
```

```
-----  
Bitmap Heap Scan on mytable  (cost=12.01..16.02 rows=1 width=64)
```

```
  Recheck Cond: (mydocb @> '{"company": "Magnafone"}'::jsonb)
```

```
    -> Bitmap Index Scan on mytable_mydocb_idx  (cost=0.00..12.01 rows=1 width=0)
```

```
        Index Cond: (mydocb @> '{"company": "Magnafone"}'::jsonb)
```

JSON Path Type

The `jsonpath` type implements support for the SQL/JSON path language.

Semantics:

- Dot (.) is used for member access.
- Square brackets ([]) are used for array access.
- SQL/JSON arrays are 0-relative, unlike regular SQL arrays that start from 1.

<https://www.postgresql.org/docs/current/datatype-json.html>

<https://www.postgresql.org/docs/current/functions-json.html>

Operators

EXAMPLE: (Operator ->)

Get JSON array element

(indexed from zero, negative integers count from the end)

```
SELECT '["a":"foo"],["b":"bar"],["c":"baz"]'::json->2 as "operator ->";
```

```
operator ->
```

```
-----
```

```
{"c":"baz"}
```

Operators: JSON and JSONB

Operator	Example	Example Result
->	'{"a":"foo"}, {"b":"bar"}, {"c":"baz"}'::json->2	{ "c": "baz" }
->	'{"a": {"b":"foo"}}'::json->'a'	{ "b": "foo" }
->>	'[1,2,3]'::json->>2	3
->>	'{"a":1, "b":2}'::json->>'b'	2
#>	'{"a": {"b":{"c": "foo"}}}'::json#>'{a,b}'	{ "c": "foo" }
#>>	'{"a":[1,2,3], "b":[4,5,6]}'::json#>>'{a,2}'	3
@>	'{"a":1, "b":2}'::jsonb @> '{"b":2}'::jsonb	t
<@	'{"b":2}'::jsonb <@ '{"a":1, "b":2}'::jsonb	t
?	'{"a":1, "b":2}'::jsonb ? 'b'	t
?	'{"a":1, "b":2, "c":3}'::jsonb ? array['b', 'c']	t
?&	'["a", "b"]'::jsonb ?& array['a', 'b']	t
	'["a", "b"]'::jsonb ['c', 'd']::jsonb	["a", "b", "c", "d"]
-	'{"a": "b"}'::jsonb - 'a'	{ }
-	'{"a": "b", "c": "d"}'::jsonb - '{a,c}'::text[]	{ }
-	'["a", "b"]'::jsonb - 1	["a"]
#-	'["a", {"b":1}]'::jsonb #- '{1,b}'	["a", { }]
@?	'{"a":[1,2,3,4,5]}'::jsonb @? '\$.a[*] ? (@ > 2)'	t
@@	'{"a":[1,2,3,4,5]}'::jsonb @@ '\$.a[*] > 2'	t

Operators: JSONB only

Operator Details

<	less than
>	greater than
<=	less than or equal to
>=	greater than or equal to
=	equal
<> or !=	not equal

Functions

```
SELECT to_json('Fred said "Hi." '::text);
       to_json
```

```
-----
"Fred said \"Hi.\""
```

```
SELECT array_to_json('{{1,5},{99,100}}'::int[]);
       array_to_json
```

```
-----
[[1,5],[99,100]]
```

```
SELECT * FROM json_each('{ "a": "foo", "b": "bar" }');
```

key	value
a	"foo"
b	"bar"

Functions Cont'd

```
row_to_json(row(1,'foo'))  
json_build_array(1,2,'3',4,5)  
json_build_object('foo',1,'bar',2)  
json_object('{a, 1, b, "def", c, 3.5}') json_object('{{a, 1},{b, "def"},{c, 3.5}}')  
json_object('{a, b}', '{1,2}')  
json_array_length('[1,2,3,{"f1":1,"f2":[5,6]},4]')
```

```
{"f1":1,"f2":"foo"}  
[1, 2, "3", 4, 5]  
{"foo": 1, "bar": 2}  
{"a": "1", "b": "def", "c": "3.5"}  
{"a": "1", "b": "2"}  
5
```

Functions Cont'd

```
json_extract_path(from_json json, VARIADIC path_elems text[]) jsonb_extract_path(from_json jsonb, VARIADIC path_elems text[])
json_extract_path_text(from_json json, VARIADIC path_elems text[]) jsonb_extract_path_text(from_json jsonb, VARIADIC path_elems text[])
json_object_keys(json) jsonb_object_keys(jsonb)
json_populate_record(base anyelement, from_json json) jsonb_populate_record(base anyelement, from_json jsonb)
json_populate_recordset(base anyelement, from_json json) jsonb_populate_recordset(base anyelement, from_json jsonb)
json_array_elements(json) jsonb_array_elements(jsonb)
json_array_elements_text(json) jsonb_array_elements_text(jsonb)
json_typeof(json) jsonb_typeof(jsonb)
json_to_record(json) jsonb_to_record(jsonb)
json_to_recordset(json) jsonb_to_recordset(jsonb)
json_strip_nulls(from_json json) jsonb_strip_nulls(from_json jsonb)
jsonb_set(target jsonb, path text[], new_value jsonb [, create_missing boolean])
jsonb_insert(target jsonb, path text[], new_value jsonb [, insert_after boolean])
```

Make Sense of Your Database

Part II: Scenarios

Create Database

- Populate Database
- Translate Relational to JSON
- Queries

Create Database

```
BEGIN;
create extension pgcrypto;
create or replace function f_password_hash()
returns trigger
as
$$
begin
    NEW.password = crypt(NEW.password::text, gen_salt('md5'));
    RETURN NEW;
end;
$$
language plpgsql;

create type type_communication as enum ('email', 'slack', 'skype');

create table if not exists identity (
    id_identity integer primary key generated by default as identity,
    first_name varchar(15) not null,
    last_name varchar(60) not null,
    ss char(9) unique not null check (ss ~ '^[[:digit:]]{9}$')
);

create table if not exists address (
    id_address integer primary key generated by default as identity,
    id_identity integer references identity(id_identity) not null,
    street_address varchar(128) null,
    city varchar(64) null,
    zip char(5) null check (zip ~ '^[[:digit:]]{5}$')
);

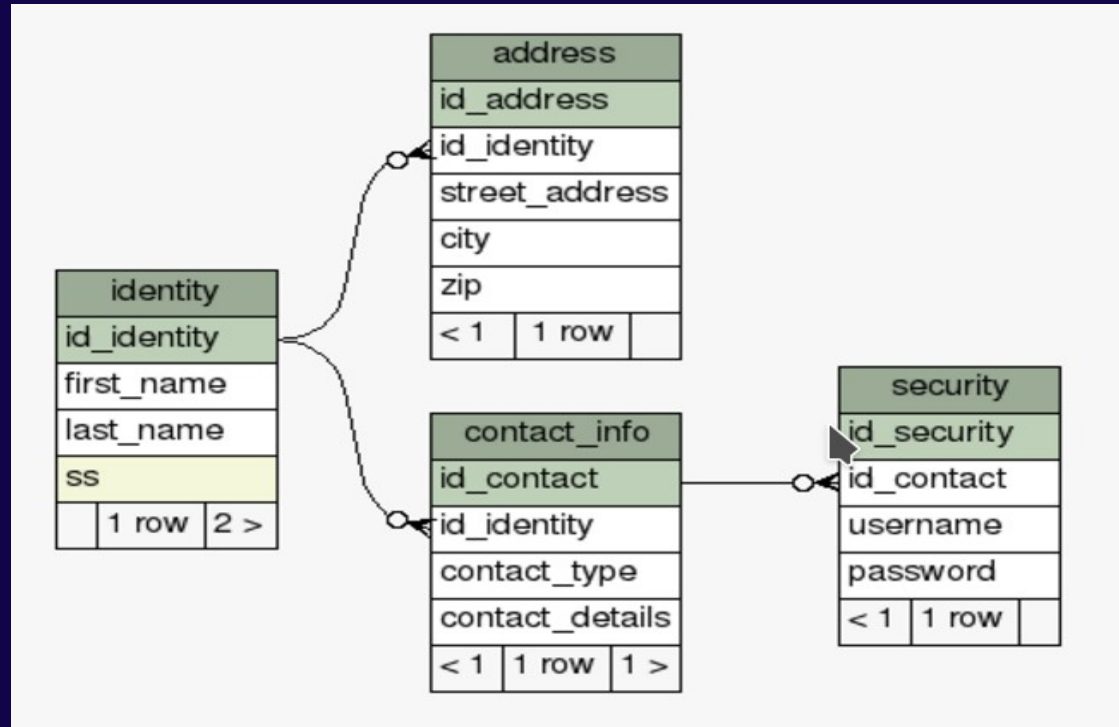
create table if not exists contact_info (
    id_contact integer primary key generated by default as identity,
    id_identity integer references identity(id_identity) not null,
    contact_type type_communication not null,
    contact_details varchar(64) not null
);
```

```
create table if not exists security (
    id_security integer primary key generated by default as identity,
    id_contact integer references contact_info (id_contact),
    username varchar(16) not null,
    password varchar not null
);

create table document (
    id_doc integer primary key generated by default as identity,
    myrecord json
);

create trigger tr_password_hash
before insert or update
on security
for each row
execute function f_password_hash();
COMMIT;
```


Entity Relationship Diagram



Populate Database

```
BEGIN;
```

```
insert into identity(first_name,last_name,ss)
  values ('robert','bernier',123456789) returning id_identity;
```

```
insert into address (id_identity, street_address, city, zip)
  values (1,'123 my address','my city','12345') returning id_address;
```

```
insert into contact_info (id_identity, contact_type, contact_details)
  values (1, 'email','robert.bernier@percona.com'),
  (1, 'skype','rbernier_zulu'),
  (1, 'slack','rbernier') returning id_contact;
```

```
insert into security (id_contact, username, password)
  values (1,'user1','mypassword'),
  (2,'user2','mypassword'),
  (3,'user3','mypassword')
  returning id_security, password;
```

```
insert into identity(first_name,last_name,ss)
  values ('conrad','black',234567891) returning id_identity;
```

```
insert into address (id_identity, street_address, city, zip)
  values (2,'1313 mocking bird lane','smallville','23451')
  returning id_address;
```

```
insert into contact_info (id_identity, contact_type, contact_details)
  values (2, 'email','user1'),
  (2, 'skype','user2') returning id_contact;
```

```
COMMIT;
ANALYZE;
```

Translate Relational to JSON

row_to_json

```
-----  
\df row_to_json  
  
                List of functions  
 Schema      | Name      | Result data type | Argument data types | Type  
-----+-----+-----+-----+-----  
 pg_catalog  | row_to_json | json              | record               | func  
 pg_catalog  | row_to_json | json              | record, boolean     | func
```

Translate Relational to JSON

EX: row_to_json

```
with a (i,j,k,l,m,n,o) as
(
  select id_identity,
         first_name,
         last_name,
         ss,
         street_address,
         city m,
         zip n
  from identity
  join address using (id_identity)
),
b as (select row_to_json(row(i,j,k,l,m,n,o)) as r from a)
select * from b;
```

r

```
{"f1":1,"f2":"robert","f3":"bernier","f4":"123456789","f5":"123 my address","f6":"my city","f7":"12345"}
{"f1":2,"f2":"conrad","f3":"black","f4":"234567891","f5":"1313 mocking bird lane","f6":"smallville","f7":"23451"}
```

Translate Relational to JSON

EX cont'd: row_to_json

with a (i,j,k) as

```
(
  select id_identity,
         username,
         password
  from security
  join contact_info using (id_contact)
  join identity using (id_identity)
)
select row_to_json(row(i,j,k)) as r from a;
```

r

```
-----
{"f1":1,"f2":"user1","f3":"$1$alvLhvAx$iseOfgiKDbvRrWRPz9Jpb/" }
{"f1":1,"f2":"user2","f3":"$1$VADRCU8H$U5YGB7NNJWM9vdSe9QJah." }
{"f1":1,"f2":"user3","f3":"$1$ko4Yb/wR$j.6/eEIcuJ.4rAEL8rpuo1" }
```

Translate Relational to JSON

EX: Building The Document

with

```
a as (select id_identity,
            username,
            password
      from security
      join contact_info using (id_contact)
      join identity using (id_identity)),
b as (select json_agg(json_build_object(username,password)) from a),
c as (select json_build_object('username_password',json_agg) as charlie from b),
d as (select json_build_object (
      'id_identity', id_identity,
      'first_name',  first_name,
      'last_name',   last_name,
      'ss',          ss,
      'street_address',street_address,
      'city',        city,
      'zip',         zip
    ) as delta from identity
      join address using (id_identity))
insert into document(myrecord)
select delta::jsonb || charlie::jsonb from c,d;
```

Queries

```
-- QUERY 1:  
select myrecord from document;
```

```
{  
  "ss": "123456789", "zip": "12345", "city": "my city", "last_name": "bernier", "first_name":  
  "robert", "id_identity": 1, "street_address": "123 my address", "username_password":  
  [{  
    "robert.bernier": "$1$LZoB/slh$W.stf.vWGrR7/diNJ/UeU1",  
    "rbernier_zulu": "$1$uGnN/7tV  
$FiFem5POZKILYnGAjNutm0",  
    "rbernier": "$1$MPPrHGpH$dqeVGLSAl2cT2CS44MTUa0",  
    "conrad.black":  
    "$1$6uDGVVu9$qIKVZht3EYjdjIIKT1YX00",  
    "cblack_literary": "$1$dibFFjBx$23/xMXavCVc0IscDPM00f0"}  
  ]  
}
```

Queries

```
-- QUERY 2:  
select myrecord -> 'first_name' as "first name" from document limit 1;  
  
first name  
-----  
"robert"
```


Queries

```
-- QUERY 3:  
select myrecord -> 'last_name' as "last name" from document limit 1;  
  
last name  
-----  
"bernier"
```

Queries

```
-- QUERY 4:  
select myrecord -> 'username_password' as "username/passsswords" from document limit 1;  
-----  
[{"robert.bernier": "$1$AKXaaSQM$0029cNu0PHb/7UdoF/CbV1"}, {"rbernier_zulu":  
"$1$arVZSMo3$UVCjAVhyhXQmfpnYDkpIY1"}, {"rbernier": "$1$Q7urUIj0$Sm992oUd7Y4dc/QruSV7Z0"}]
```

Queries

```
-- QUERY 5:  
select myrecord -> 'username_password' ->> 0 as "username/password" from document limit 1;  
  
username/password  
-----  
{"robert.bernier": "$1$AKXaaSQM$0029cNu0PHb/7UdoF/CbV1"}
```

Queries

```
-- QUERY 6:  
select myrecord -> 'username_password' ->> 1 as "username/password" from document limit 1;  
  
username/password  
-----  
{"rbernier_zulu": "$1$arVZSMo3$UVCjAVhyhXQmfpnYDkpIY1"}
```

Queries

```
-- QUERY 7:
select myrecord -> 'username_password' ->> 2 as "username/password" from document limit 1;

| | | username/password
-----
{"rbernier": "$1$Q7urUIj0$Sm992oUd7Y4dc/QruSV7Z0"}
```

Queries

```
-- QUERY 8:  
select (myrecord -> 'username_password' ->> 0)::json -> 'robert.bernier' as "my password" from  
document limit 1;  
  
-----  
my password  
  
"$1$AKXaaSQM$0029cNu0PHb/7UdoF/CbV1"
```

Queries

```
-- QUERY 9:  
select myrecord -> 'first_name' as "first name",  
        myrecord -> 'last_name' as "last name",  
        myrecord -> 'street_address' as "street address",  
        myrecord -> 'city' as "city",  
        myrecord -> 'zip' as "zip"  
from document;
```

first name	last name	street address	city	zip
"robert"	"bernier"	"123 my address"	"my city"	"12345"
"conrad"	"black"	"1313 mocking bird lane"	"smallville"	"23451"

Queries

```
-- QUERY 10:
select myrecord -> 'first_name' as "first name",
       myrecord -> 'last_name' as "last name",
       myrecord -> 'street_address' as "street address",
       myrecord -> 'city' as "city",
       myrecord -> 'zip' as "zip"
from document
where
  myrecord ->> 'first_name' = 'conrad'
and myrecord ->> 'last_name' = 'black';
```

first name	last name	street address	city	zip
"conrad"	"black"	"1313 mocking bird lane"	"smallville"	"23451"

Questions?

Thank You!

Thank you

Message for the end of the presentation goes here

Name goes here

 Contact/social media

 Contact/social media

 Contact/social media